

Dataset: Haberman's Survival Data

Link: <https://www.kaggle.com/gilsousa/habermans-survival-data-set>
(<https://www.kaggle.com/gilsousa/habermans-survival-data-set>)

Relevant Information: The dataset contains cases from a study that was conducted between 1958 and 1970 at the University of Chicago's Billings Hospital on the survival of patients who had undergone surgery for breast cancer.

Objective: Classify as to which patient survived for more than 5 years and who died within 5 years.

Attribute Information:

- 1) Age - Age of patient at time of operation
- 2) Year - Patient's year of operation
- 3) Axl_nodes - Number of positive axillary nodes detected
- 4) Survival - Survival status (class attribute) 1 = the patient survived 5 years or longer, 2 = the patient died within 5 year

```
In [31]: #Importing the necessary Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [32]: #Reading the data
df = pd.read_csv("haberman.csv", names=["Age", "Year", "Axl_nodes", "Survival"])
```

```
In [33]: df.head()
```

```
Out[33]:
```

	Age	Year	Axl_nodes	Survival
0	30	64	1	1
1	30	62	3	1
2	30	65	0	1
3	31	59	2	1
4	31	65	4	1

```
In [34]: df.shape
```

```
Out[34]: (306, 4)
```

In [35]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 306 entries, 0 to 305
Data columns (total 4 columns):
Age          306 non-null int64
Year         306 non-null int64
Axl_nodes    306 non-null int64
Survival     306 non-null int64
dtypes: int64(4)
memory usage: 9.6 KB
```

In [36]: `df.describe()`

Out[36]:

	Age	Year	Axl_nodes	Survival
count	306.000000	306.000000	306.000000	306.000000
mean	52.457516	62.852941	4.026144	1.264706
std	10.803452	3.249405	7.189654	0.441899
min	30.000000	58.000000	0.000000	1.000000
25%	44.000000	60.000000	0.000000	1.000000
50%	52.000000	63.000000	1.000000	1.000000
75%	60.750000	65.750000	4.000000	2.000000
max	83.000000	69.000000	52.000000	2.000000

Observation

1) It can be seen that youngsters and children of the age group 0 to 29 are not present in dataset and may be less likely prone to cancer.

In [37]: `df["Survival"].value_counts()`

```
Out[37]: 1    225
         2     81
         Name: Survival, dtype: int64
```


Observation

1) Seeing the Survival Value_counts(), it can be said that Survival column is imbalanced.

UNI-VARIATE ANALYSIS

In [38]: *#EDA of Age*

```
sns.set_style("whitegrid")
sns.distplot(a = df["Age"], bins = 10, kde = False)
plt.xlabel("Age in years")
plt.ylabel("Count")
plt.title("Age Distribution")
plt.show()
```

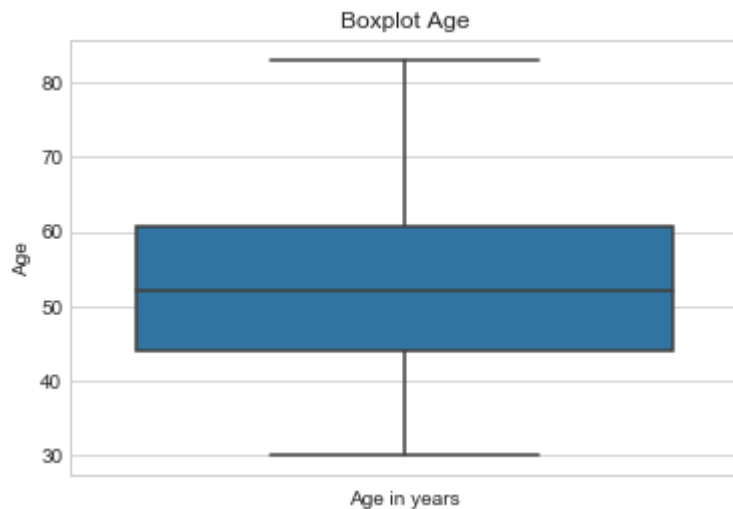


In [39]: *#Age Boxplot*

```
sns.set_style("whitegrid")
sns.boxplot(x = "Age", data = df, orient = "v")
plt.xlabel("Age in years")
plt.title("Boxplot Age")
plt.show()

age_middle = len(df[(df["Age"]>=40) & (df["Age"]<=60)]) / len(df["Age"])
age_young = len(df[(df["Age"]>=0) & (df["Age"]<=39)]) / len(df["Age"])
age_old = len(df[df["Age"]>=61]) / len(df["Age"])

print("Percentage of people in age group of 0 to 39 are:", age_young*100, "%")
print("Percentage of people in age group of 40 to 60 are:", age_middle*100, "%")
print("Percentage of people in age group greater than 60 are:", age_old*100, "%")
```



Percentage of people in age group of 0 to 39 are: 13.071895424836603 %
Percentage of people in age group of 40 to 60 are: 61.76470588235294 %
Percentage of people in age group greater than 60 are: 25.163398692810457 %

```

In [40]: #PDF
#Calculating PDF of Age Feature


count, bin_edges = np.histogram(a = df["Age"], bins =10)
pdf = count/sum(count)

#Calculating CDF
a = 0
cdf = []
for i in count:
    a = a+i
    cdf.append(a)

z = cdf/sum(count)

plt.plot(bin_edges[1:],pdf)
plt.plot(bin_edges[1:],z)
plt.xlabel("Age")
plt.ylabel("Probablity")
plt.title("PDF/CDF vs Age in years")
legend = ["PDF","CDF"]
plt.legend(legend)
plt.show()

```



Observation:

- 1) The Major chunk of people (around 62%) who are prone to cancer belongs to middle age group(40 to 60 years).
- 2) Even in CDF we can see that 80 percentile of people have an age of 62 or less.

In []:


In [41]: *#EDA of Axillary Nodes*

```
sns.distplot(df["Ax1_nodes"], kde = False)
plt.xlabel("Axillary Nodes")
plt.ylabel("Count")
plt.title("Axillary node distribution")
plt.show()

count_0 = len(df[df["Ax1_nodes"]==0]["Ax1_nodes"])
count_1 = len(df[df["Ax1_nodes"]==1]["Ax1_nodes"])
count_2 = len(df[df["Ax1_nodes"]==2]["Ax1_nodes"])

tot = len(df)

print("Percentage of Axillary Node between 0 to 2: ",((count_0+count_1+count_2)/tot)*100,"%")
```



Percentage of Axillary Node between 0 to 2: 64.37908496732027 %

```

In [42]: #PDF
#Calculating PDF of Axillary Node


count, bin_edges = np.histogram(a = df["Ax1_nodes"], bins =50)
pdf = count/sum(count)

#Calculating CDF
a = 0
cdf = []
for i in count:
    a = a+i
    cdf.append(a)

z = cdf/sum(count)

plt.plot(bin_edges[1:],pdf)
plt.plot(bin_edges[1:],z)
plt.xlabel("Axillary Nodes")
plt.ylabel("Probablity")
plt.title("PDF/CDF of Axillary Nodes")
legend = ["PDF","CDF"]
plt.legend(legend)
plt.show()

```



Observation:

- 1) More than 60 percent of the Axillary nodes are between 0 and 2.
- 2) PDF for Axillary nodes is highly left skewed/Negatively skewed, which means most of the Axillary Nodes lie between 0 to 10.

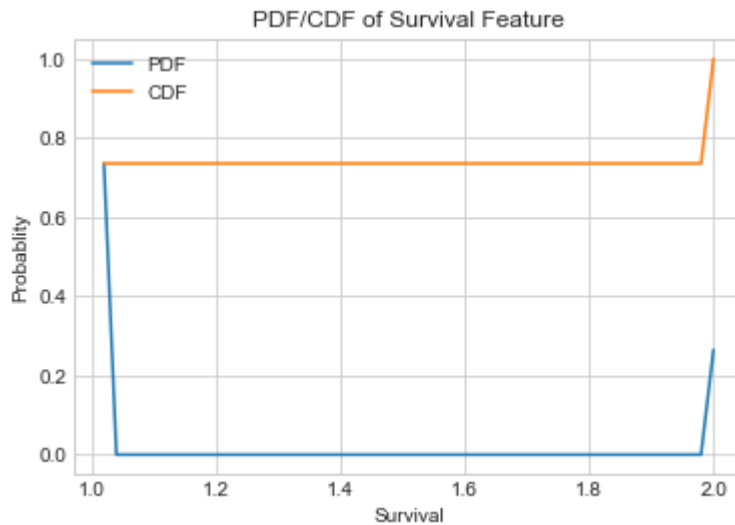
```
In [43]: #PDF
#Calculating PDF of Survival

count, bin_edges = np.histogram(a = df["Survival"], bins =50)
pdf = count/sum(count)

#Calculating CDF
a = 0
cdf = []
for i in count:
    a = a+i
    cdf.append(a)

z = cdf/sum(count)

plt.plot(bin_edges[1:],pdf)
plt.plot(bin_edges[1:],z)
plt.xlabel("Survival")
plt.ylabel("Probablity")
plt.title("PDF/CDF of Survival Feature")
legend = ["PDF","CDF"]
plt.legend(legend)
plt.show()
```




```

In [44]: #PDF
#Calculating PDF of Year

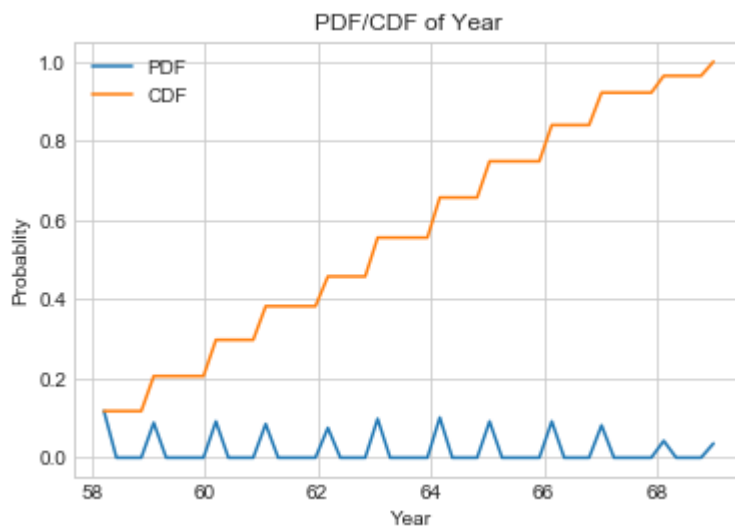
count, bin_edges = np.histogram(a = df["Year"], bins =50)
pdf = count/sum(count)

#Calculating CDF
a = 0
cdf = []
for i in count:
    a = a+i
    cdf.append(a)

z = cdf/sum(count)

plt.plot(bin_edges[1:],pdf)
plt.plot(bin_edges[1:],z)
plt.xlabel("Year")
plt.ylabel("Probablity")
plt.title("PDF/CDF of Year")
legend = ["PDF","CDF"]
plt.legend(legend)
plt.show()

```



Observation

1) It can be seen that the number of cancer cases for all years mentioned have been almost constant except for year 1968.

BI-Variate Analysis


In [45]: *#Age vs Axillary Node*

```
ax = sns.lmplot(x="Ax1_nodes", y="Age", data=df, fit_reg=False, hue="Survival")
plt.xlabel("Axillary nodes")
plt.ylabel("Age")
plt.title("Age vs Axillary nodes")

#handles, labels = ax.get_legend_handles_labels()
#plt.legend(handles, ["Survived>5 years", "Survived<5 years"])
plt.show()
```



```
In [46]: ax = sns.boxplot(data = df, x="Survival", y="Age", hue = "Survival")
plt.xlabel("Survival")
plt.ylabel("Age in years")
plt.title("Age distribution according to Survival Rate")
handles, labels = ax.get_legend_handles_labels()
plt.legend(handles,["Survived>5 yaers", "Survived<5 years"])
plt.show()
```

**Observation:**

1) Age cannot be considered a major factor for Survival in this case because many people of similar ages have survived as well as died due to the disease.

In [47]: `#Pairplot`

```
ax = sns.pairplot(data = df, hue = "Survival", vars = ["Age", "Year", "Ax1_nodes"])
plt.title("Pairplot")
plt.show()
```



Observation

- 1) Age seems to be normally distributed.
- 2) For most of the people, the Axillary node seems to be in range of 0 to 10.

In [48]: *#Boxplot - Auxillary Node vs Survival*

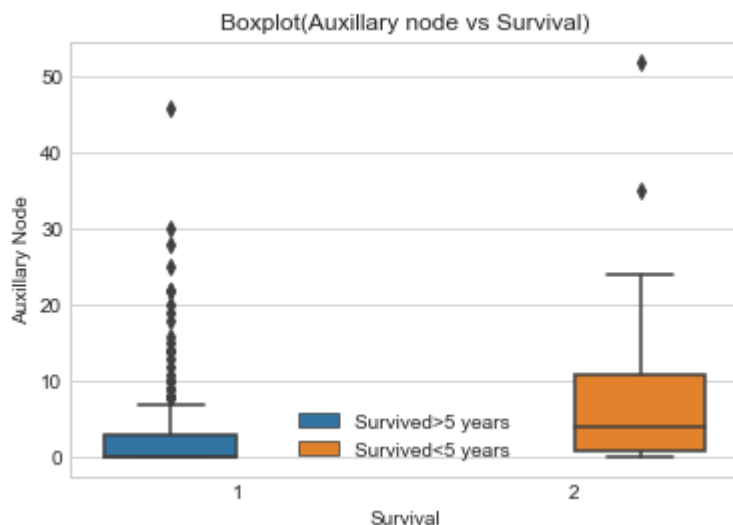
```
ax = sns.boxplot(data = df, x="Survival", y="Ax1_nodes", hue="Survival")
plt.title("Boxplot(Auxillary node vs Survival)")
plt.xlabel("Survival")
plt.ylabel("Auxillary Node")
handles, labels = ax.get_legend_handles_labels()
plt.legend(handles, ["Survived>5 years", "Survived<5 years"])
plt.show()

count_0 = len(df[(df["Survival"]==1) & (df["Ax1_nodes"]==0)])
count_1 = len(df[(df["Survival"]==1) & (df["Ax1_nodes"]==1)])
count_2 = len(df[(df["Survival"]==1) & (df["Ax1_nodes"]==2)])
count_3 = len(df[(df["Survival"]==1) & (df["Ax1_nodes"]==3)])
tot = len(df[df["Survival"]==1])

print("Percentage of People who Survived more than 5 years with Auxillary nodes between 0 and 3:", ((count_0+count_1+count_2+count_3)/tot)*100,"%")

count = []
for i in range(0,12):
    count.append(len(df[(df["Survival"]==2) & (df["Ax1_nodes"]==i)]))
tot = len(df[df["Survival"]==2])

print("Percentage of people who could not survive more than 5 years with an Auxillary node between 3 and 11:",(sum(count)/tot)*100,"%")
```



Percentage of People who Survived more than 5 years with Auxillary nodes between 0 and 3: 79.11111111111111 %

Percentage of people who could not survive more than 5 years with an Auxillary node between 3 and 11: 75.30864197530865 %

Observation:

- 1) About 79 percent of people who survived for more than 5 years had an Axillary node between 0 and 3.
- 2) 75 percent of people who could not survive for more than 5 years had Axillary node ranging from 3 to 11.
- 3) Many outliers can be seen in the survival plot whose axillary node ranges from 7 to 32.

In [49]: *#Correlation Table*

```
df.corr(method = "spearman")
```

Out[49]:

	Age	Year	Axl_nodes	Survival
Age	1.000000	0.091069	-0.097691	0.055914
Year	0.091069	1.000000	-0.036146	-0.007536
Axl_nodes	-0.097691	-0.036146	1.000000	0.327081
Survival	0.055914	-0.007536	0.327081	1.000000

Observation:

- 1) Looking at the correlationship of features with eachother, it can be said that "Axillary nodes" and "Survival" are most related.

In [50]: `df[df["Survival"]==1]["Axl_nodes"].describe()`


```
Out[50]: count    225.000000
mean         2.791111
std          5.870318
min          0.000000
25%          0.000000
50%          0.000000
75%          3.000000
max         46.000000
Name: Axl_nodes, dtype: float64
```

In [51]: `df[df["Survival"]==2]["Axl_nodes"].describe()`

```
Out[51]: count    81.000000
mean         7.456790
std          9.185654
min          0.000000
25%          1.000000
50%          4.000000
75%         11.000000
max         52.000000
Name: Axl_nodes, dtype: float64
```


In [52]: *#EDA Violinplot*

```
ax = sns.violinplot(x = "Survival", y = "Ax1_nodes", hue= "Survival", data = df
)
plt.xlabel("Survival")
plt.ylabel("Axillary Nodes")
plt.title("Violinplot(Distribution of Axillary nodes of the Survived People)")
handles, labes = ax.get_legend_handles_labels()
plt.legend(handles, ["Survived>5 years","Survived<5 years"])
plt.show()
```



In [53]: *#EDA Count of Axillary node for Survival = 1*


```
sns.distplot(df[df["Survival"]==1]["Ax1_nodes"], kde = False)
plt.xlabel("Axillary node")
plt.ylabel("Count of the nodes")
plt.title("Distribution of Axillary nodes for Survival = 1")
plt.show()
```



Observation

1) It can be seen that the graph is highly left skewed and can be seen that there are a lot of outliers.

```
In [54]: sns.distplot(df[df["Survival"]==2]["Ax1_nodes"], kde = False)
plt.xlabel("Axillary node")
plt.ylabel("Count of the nodes")
plt.title("Distribution of Axillary nodes for Survival = 2")
plt.show()
```



Observation

1) It can be seen that the count of Axillary node is less left skewed for people who survived for < 5 years. It also has less number of outliers.

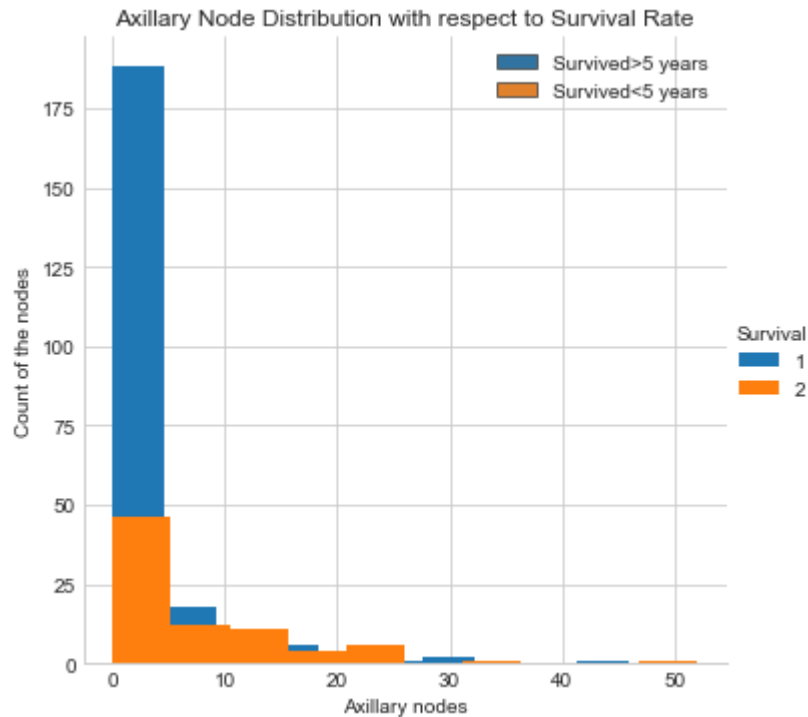
In []:

```
In [55]: df["Ax1_nodes"].value_counts().head()
```

```
Out[55]: 0    136
         1     41
         2     20
         3     20
         4     13
         Name: Ax1_nodes, dtype: int64
```



```
In [56]: sns.FacetGrid(df, hue="Survival", size=5).map(plt.hist, "Ax1_nodes").add_legen
d();
plt.xlabel("Axillary nodes")
plt.ylabel("Count of the nodes")
plt.title("Axillary Node Distribution with respect to Survival Rate")
handles, labels = ax.get_legend_handles_labels()
plt.legend(handles, ["Survived>5 years", "Survived<5 years"])
plt.show();
```



Observation:

1) People having Axillary node between 0-5 and survived for < 5 years are less than 1/3 rd population of people who survived, as seen in the graph.

In []:

Conclusion:

- 1) Young people and children are less prone to Cancer.
- 2) The label column (Survival) is highly imbalanced and hence some measures need to be taken to convert it to a balanced column.
- 3) The major chunk of population are middle aged people(40-60) with 61 % of population. Though Age doesn't play a major part to describe the fact who survived for less than 5 years or greater than 5 years.
- 4) A high correlation can be seen between Axillary Node and Survival Rate, which also makes Axillary Node the best available feature for predicting the Survival Rate.
- 5) Axillary Node feature is also imbalanced with 64% of nodes ranging between 0 and 2.
- 6) People having Axillary node between 0-5 and survived for < 5 years are less than 1/3 rd population of people who survived.

In []: